

course-work

April 29, 2023

1 Course work

0. Import the libraries that you will need

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

1. Get the data - in the cell below run: **Note** you only need to run this command the first time you do the exercise. If you save and go away and come back, then can skip straight to step 2.

```
!python get-my-data.py
```

2. Read in the csv:

```
df = pd.read_csv('coursework-data.csv')
```

3. Perform some exploratory data analysis to clean up the dataset. The code needed for this part is found in the first set of exercises that you did.
 - Remove outliers
 - If any pairs of variables are highly correlated, remove one of the pair - highly correlated in this case > 0.99
4. Fit a baseline model, linear regression to map the control parameters (all parameters *except* `gllbsc_gap`) to the dependent parameter `gllbsc_gap`. Summarise its performance.

To set up the data use:

```
x = df.loc[:, df.columns != "gllbsc_gap"].values
y = df.loc[:, df.columns == "gllbsc_gap"].values
```

The rest of the code you need for this found in the second set of exercises that you did.

- From looking at the linear regression model, which features have the greatest influence on the band gap?
5. Develop a gradient boosted regressor to the same data. Summarise its performance.

1.1 Important notes

1.1.1 Submitting the coursework

When you are finished with the coursework - use **File > Save and Export Notebook As > pdf** to download a pdf of the completed notebook. Submit this pdf *via* the portal on QMplus.

The deadline for submission is Friday 28th April at 16:00.

1.1.2 Text explanations

Please please please add text to explain what you are doing in the code. Adding text boxes is easy, just add a new cell as normal then change the type to **Markdown** with the dropdown menu at the top of the cells. Adding text will make sure that markers can give you proper grades even if you make a small slip in your code. If you have no text explanation and still have a small slip, you will likely get no marks!

1.1.3 Datasets

All of your datasets are generated randomly. So do not expect the same answers as your friends. If you compare answers and find that you have something very different, do not worry.

1.1.4 Warnings from the code

Don't worry if the code throws some warnings sometimes. If it keeps running then it is fine. Warnings usually just alert you to future planned changes in the code you are using.

1.1.5 Long run times

There is a certain part of the exercise where a grid search is required. It could take quite a long time with this code. I have tested it and it took about 15 minutes for a 10-fold cross validation on a 5x5 gridsearch. Dont worry if it seems to be running for a long time, that's okay.

```
[1]: #data manipulation
import pandas as pd
import numpy as np

#data visualization
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import rcParams
```

```
[2]: #Getting the data
!python get-my-data.py
```

File already exists no more to see here.

```
[3]: #Reading in the csv file
df = pd.read_csv('coursework-data.csv')
```

```
[4]: #Getting the shape of the character in terms of (examples, characteristics)
df.shape
```

```
[4]: (749, 16)
```

```
[6]: #Using describe to get a summary of the data
df.describe()
```

```
[6]:
```

	GS mean	GS dev	HOMO_energy	Weight dev	Eneg dev	\
count	749.000000	749.000000	749.000000	749.000000	749.000000	
mean	20.129941	145.794327	-0.318680	168.245383	0.880011	
std	8.758476	3653.481944	0.041349	3652.708662	0.202963	
min	8.332000	0.000000	-0.338381	0.000000	0.000000	
25%	14.012447	6.064600	-0.338381	18.529332	0.759008	
50%	17.530000	9.299354	-0.338381	32.680587	0.899592	
75%	23.485833	16.631250	-0.320380	48.093618	1.036694	
max	80.211667	100000.000000	-0.144272	100000.000000	1.325000	

	Spg dev	gllbsc_gap	NValence mean	LUMO values	CovRad dev	\
count	749.000000	749.000000	749.000000	749.000000	749.000000	
mean	217.080871	5.308134	6.900038	481.708041	41.414244	
std	3650.926077	2.317242	2.628995	223.128903	12.558078	
min	0.000000	0.144493	2.666667	54.800000	0.000000	
25%	72.592593	3.652201	4.923077	290.248000	32.395062	
50%	89.306122	5.346850	6.300000	445.613333	40.592593	
75%	99.750000	6.903045	8.000000	656.466667	50.000000	
max	100000.000000	11.560321	20.000000	1220.600000	89.000000	

	MeltT mean	Number dev	Periodic nature	Mendeleev dev	NdValence dev	\
count	749.000000	749.000000	749.000000	749.000000	749.000000	
mean	481.708041	14.144738	-0.318680	22.488916	2.220428	
std	223.128903	7.934567	0.041349	12.110542	1.677656	
min	54.800000	0.000000	-0.338381	0.000000	0.000000	
25%	290.248000	7.836735	-0.338381	11.555556	0.555556	
50%	445.613333	13.500000	-0.338381	24.612245	2.370370	
75%	656.466667	19.222222	-0.320380	32.520000	3.750000	
max	1220.600000	36.122449	-0.144272	44.571429	5.000000	

	MeltT dev
count	749.000000
mean	493.400798
std	285.198279
min	0.000000
25%	255.460408
50%	439.656198
75%	708.573750
max	1372.535000

```
[7]: #Use info to get full list of characteristics and their data types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 749 entries, 0 to 748
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
#   ...
```

```

---  -----
0  GS mean      749 non-null    float64
1  GS dev       749 non-null    float64
2  HOMO_energy  749 non-null    float64
3  Weight dev   749 non-null    float64
4  Eneg dev     749 non-null    float64
5  Spg dev      749 non-null    float64
6  gllbasc_gap  749 non-null    float64
7  NValence mean 749 non-null    float64
8  LUMO values  749 non-null    float64
9  CovRad dev   749 non-null    float64
10 MeltT mean   749 non-null    float64
11 Number dev   749 non-null    float64
12 Periodic nature 749 non-null    float64
13 Mendeleev dev 749 non-null    float64
14 NdValence dev 749 non-null    float64
15 MeltT dev    749 non-null    float64
dtypes: float64(16)
memory usage: 93.8 KB

```

```

[8]: ##Starting to clean the dataset
      #Checking for duplicated rows
      df.duplicated().sum()

```

[8]: 0

```

[9]: df.nunique()

```

```

[9]: GS mean      748
      GS dev      748
      HOMO_energy   30
      Weight dev   748
      Eneg dev     665
      Spg dev      407
      gllbasc_gap   749
      NValence mean 184
      LUMO values   748
      CovRad dev    613
      MeltT mean    748
      Number dev    574
      Periodic nature  30
      Mendeleev dev  624
      NdValence dev  188
      MeltT dev     748
      dtype: int64

```

```
[10]: #Working to find and remove the Outliers
df.head()
```

```
[10]:      GS mean      GS dev  HOMO_energy  Weight dev  Eneg dev      Spg dev  \
0  20.010909  12.915702   -0.160771   18.519174   1.090909  105.619835
1  10.555455   2.109752   -0.338381   51.543246   0.643967   72.198347
2  11.680714   3.679592   -0.338381   67.325910   0.791837   88.571429
3  22.894091  17.092314   -0.273634   48.630545   1.029421  102.545455
4  30.719167  21.614167   -0.338381   36.772863   1.002500   92.750000

      gllbsc_gap  NValence mean  LUMO values  CovRad dev  MeltT mean  Number dev  \
0    3.178134      4.909091   601.531818   50.876033   601.531818    8.528926
1    6.064334      9.272727   731.734545   29.884298   731.734545   20.628099
2    5.143263      9.714286   979.142857   42.448980   979.142857   26.530612
3    3.833288      7.090909   555.856364   53.057851   555.856364   19.834711
4    3.881077      7.000000   337.477500   52.500000   337.477500   14.750000

      Periodic nature  Mendeleev dev  NdValence dev  MeltT dev
0         -0.160771      35.371901      1.652893   596.434711
1         -0.338381      11.107438      2.644628   991.732893
2         -0.338381      15.918367      1.224490  1320.489796
3         -0.273634      29.305785      2.826446   546.606942
4         -0.338381      31.250000      3.750000   283.151250
```

```
[11]: df.tail()
```

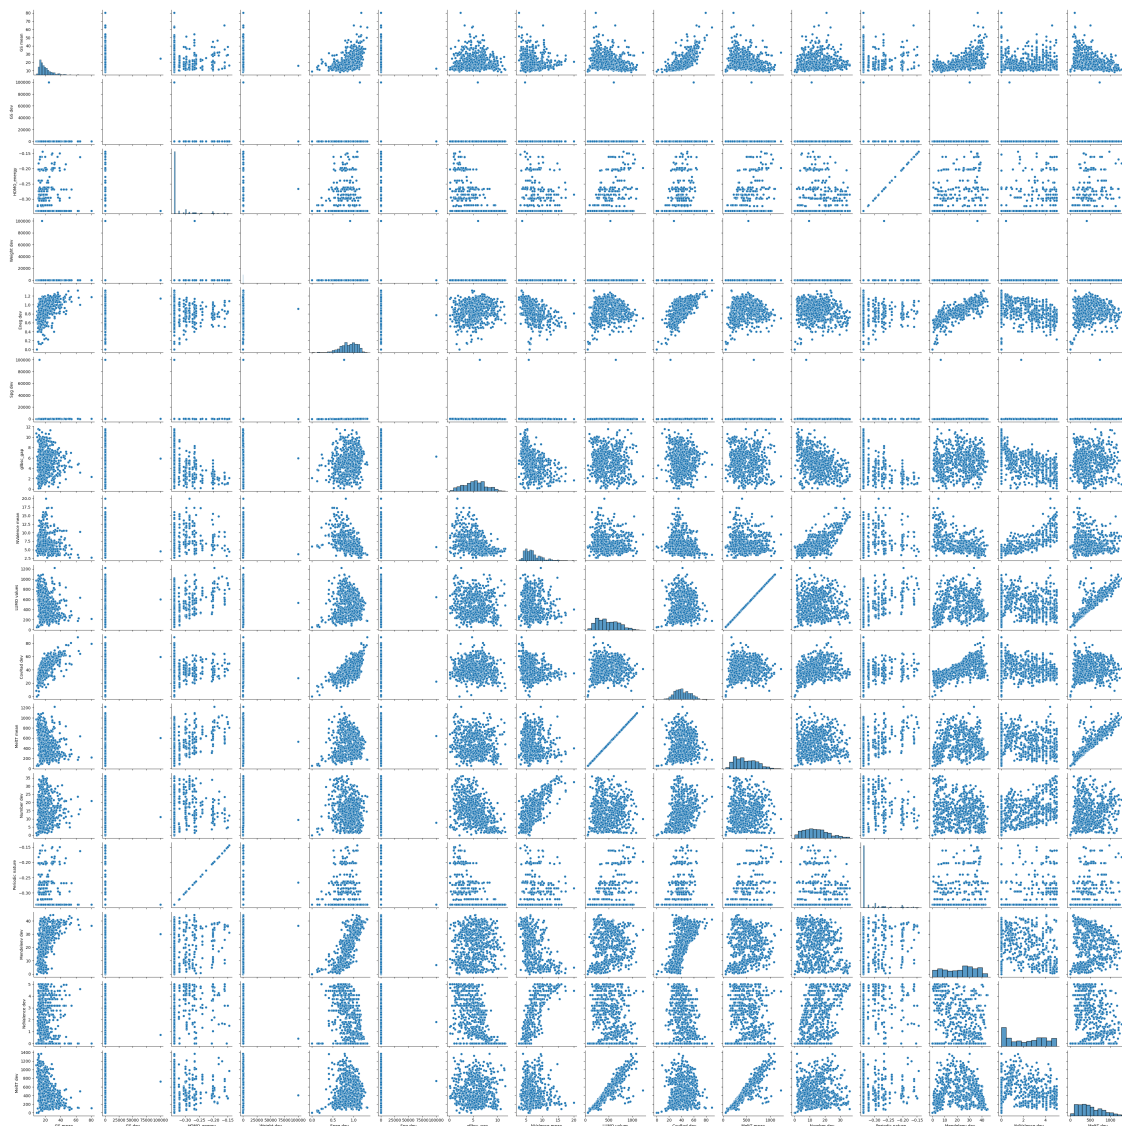
```
[11]:      GS mean      GS dev  HOMO_energy  Weight dev  Eneg dev      Spg dev  \
744  23.689231  18.415740   -0.338381   45.616776   0.962840   90.792899
745  17.307813  12.787344   -0.338381   34.121112   0.814062   82.031250
746  30.929063  24.495703   -0.197497   53.511675   0.973750   89.906250
747  14.552500   8.815417   -0.338381   18.375224   0.875972   92.458333
748  19.864250  15.972725   -0.338381   41.178064   0.988800   96.600000

      gllbsc_gap  NValence mean  LUMO values  CovRad dev  MeltT mean  \
744    7.162802      6.461538   432.184615   52.35503   432.184615
745    6.588943      6.250000   589.000000   36.75000   589.000000
746    3.423648      5.250000   825.300000   63.18750   825.300000
747    8.772669      4.666667   662.211667   30.87500   662.211667
748    4.819665      6.900000   811.659500   49.53000   811.659500

      Number dev  Periodic nature  Mendeleev dev  NdValence dev  MeltT dev
744    18.366864         -0.338381      27.479290      2.60355   464.473373
745    13.875000         -0.338381      17.250000      2.18750   667.750000
746    21.500000         -0.197497      34.875000      2.18750   577.875000
747     7.500000         -0.338381      13.222222      0.00000   708.646944
748    16.170000         -0.338381      24.930000      0.76500  1050.770250
```

```
[12]: #Graphical examination - Using seaborn to plot scatter graphs of the different
      ↪ variables against each other
      sns.pairplot(df)
```

```
[12]: <seaborn.axisgrid.PairGrid at 0x7fb09a916170>
```




```
[13]: df.columns
```

```
[13]: Index(['GS mean', 'GS dev', 'HOMO_energy', 'Weight dev', 'Eneg dev', 'Spg dev',
        'gllbsc_gap', 'NValence mean', 'LUMO values', 'CovRad dev',
        'MeltT mean', 'Number dev', 'Periodic nature', 'Mendeleev dev',
        'NdValence dev', 'MeltT dev'],
        dtype='object')
```

[14]: *#Plotting box plots to look for outliers in the columns*

```
columns = ['GS mean', 'GS dev', 'HOMO_energy', 'Weight dev', 'Eneg dev', 'Spg dev',
           'gllbsc_gap', 'NValence mean', 'LUMO values', 'CovRad dev', 'MeltT mean', 'Number dev', 'Periodic nature', 'Mendeleev dev',
           'NdValence dev', 'MeltT dev']
df[columns].plot(kind='box',subplots=True,layout=(4,4),figsize=(20,15))
plt.show()
```



[15]: *#Inspecting the individual distributions by finding the Skewness and Kurtosis of each column*

```
column_names = list(df.columns)
print(column_names)

skewness = df[columns].skew()
print(skewness)
```

```
['GS mean', 'GS dev', 'HOMO_energy', 'Weight dev', 'Eneg dev', 'Spg dev',
'gllbsc_gap', 'NValence mean', 'LUMO values', 'CovRad dev', 'MeltT mean',
'Number dev', 'Periodic nature', 'Mendeleev dev', 'NdValence dev', 'MeltT dev']
```

```

GS mean          1.978349
GS dev           27.367619
HOMO_energy      2.313798
Weight dev       27.366534
Eneg dev         -0.677817
Spg dev          27.366438
gllbsc_gap       0.097489
NValence mean    1.334078
LUMO values      0.430720
CovRad dev       0.217036
MeltT mean       0.430720
Number dev       0.495327
Periodic nature  2.313798
Mendeleev dev    -0.177724
NdValence dev    0.075290
MeltT dev        0.628195
dtype: float64

```

```
[16]: kurtosis = df[columns].kurt()
      print(skewness)
```

```

GS mean          1.978349
GS dev           27.367619
HOMO_energy      2.313798
Weight dev       27.366534
Eneg dev         -0.677817
Spg dev          27.366438
gllbsc_gap       0.097489
NValence mean    1.334078
LUMO values      0.430720
CovRad dev       0.217036
MeltT mean       0.430720
Number dev       0.495327
Periodic nature  2.313798
Mendeleev dev    -0.177724
NdValence dev    0.075290
MeltT dev        0.628195
dtype: float64

```

```
[39]: #Based on the values of Skewness and Kurtosis, it seems like the following
      ↪columns have outliers: GS dev, Weight dev, and Spg dev
      #We know this because they are the only columns with high values, of 27, for
      ↪Skewness and Kurtosis
```

```
[17]: #Repeating the describe function to find the mean of each column
      df.describe()
```



```
[17]:
```

	GS mean	GS dev	HOMO_energy	Weight dev	Eneg dev	\
count	749.000000	749.000000	749.000000	749.000000	749.000000	
mean	20.129941	145.794327	-0.318680	168.245383	0.880011	
std	8.758476	3653.481944	0.041349	3652.708662	0.202963	
min	8.332000	0.000000	-0.338381	0.000000	0.000000	
25%	14.012447	6.064600	-0.338381	18.529332	0.759008	
50%	17.530000	9.299354	-0.338381	32.680587	0.899592	
75%	23.485833	16.631250	-0.320380	48.093618	1.036694	
max	80.211667	100000.000000	-0.144272	100000.000000	1.325000	

	Spg dev	gllbsc_gap	NValence mean	LUMO values	CovRad dev	\
count	749.000000	749.000000	749.000000	749.000000	749.000000	
mean	217.080871	5.308134	6.900038	481.708041	41.414244	
std	3650.926077	2.317242	2.628995	223.128903	12.558078	
min	0.000000	0.144493	2.666667	54.800000	0.000000	
25%	72.592593	3.652201	4.923077	290.248000	32.395062	
50%	89.306122	5.346850	6.300000	445.613333	40.592593	
75%	99.750000	6.903045	8.000000	656.466667	50.000000	
max	100000.000000	11.560321	20.000000	1220.600000	89.000000	

	MeltT mean	Number dev	Periodic nature	Mendeleev dev	NdValence dev	\
count	749.000000	749.000000	749.000000	749.000000	749.000000	
mean	481.708041	14.144738	-0.318680	22.488916	2.220428	
std	223.128903	7.934567	0.041349	12.110542	1.677656	
min	54.800000	0.000000	-0.338381	0.000000	0.000000	
25%	290.248000	7.836735	-0.338381	11.555556	0.555556	
50%	445.613333	13.500000	-0.338381	24.612245	2.370370	
75%	656.466667	19.222222	-0.320380	32.520000	3.750000	
max	1220.600000	36.122449	-0.144272	44.571429	5.000000	

	MeltT dev
count	749.000000
mean	493.400798
std	285.198279
min	0.000000
25%	255.460408
50%	439.656198
75%	708.573750
max	1372.535000

```
[18]: #Using to mean, I can find out which specific row has an outlier in each
      ↪outlier column mentioned above
      df[df['GS dev'] > 145]
```

```
[18]:
```

	GS mean	GS dev	HOMO_energy	Weight dev	Eneg dev	Spg dev	\
320	24.650556	100000.0	-0.338381	25.497611	1.14449	96.761905	

	gllbsc_gap	NValence mean	LUMO values	CovRad dev	MeltT mean	\
320	5.880432	4.571429	602.081905	59.482993	602.081905	

	Number dev	Periodic nature	Mendeleev dev	NdValence dev	MeltT dev
320	11.102041	-0.338381	29.986395	0.725624	726.627664

```
[19]: #Dropping the row with the outlier from GS dev
df = df.drop(320)
```

```
[20]: df[df['Weight dev'] > 168]
```

```
[20]:
```

	GS mean	GS dev	HOMO_energy	Weight dev	Eneg dev	Spg dev	\
67	15.888457	1.166008	-0.266297	100000.0	0.90963	24.54321	

	gllbsc_gap	NValence mean	LUMO values	CovRad dev	MeltT mean	\
67	5.962006	3.703704	533.268148	27.654321	533.268148	

	Number dev	Periodic nature	Mendeleev dev	NdValence dev	MeltT dev
67	9.399177	-0.266297	36.345679	0.411523	408.404719

```
[21]: df = df.drop(67)
```

```
[22]: df[df['Spg dev'] > 217]
```

```
[22]:
```

	GS mean	GS dev	HOMO_energy	Weight dev	Eneg dev	Spg dev	gllbsc_gap	\
16	12.515	6.229	-0.338381	18.529332	0.7692	100000.0	6.291999	

	NValence mean	LUMO values	CovRad dev	MeltT mean	Number dev	\
16	5.8	645.288	22.24	645.288	7.6	

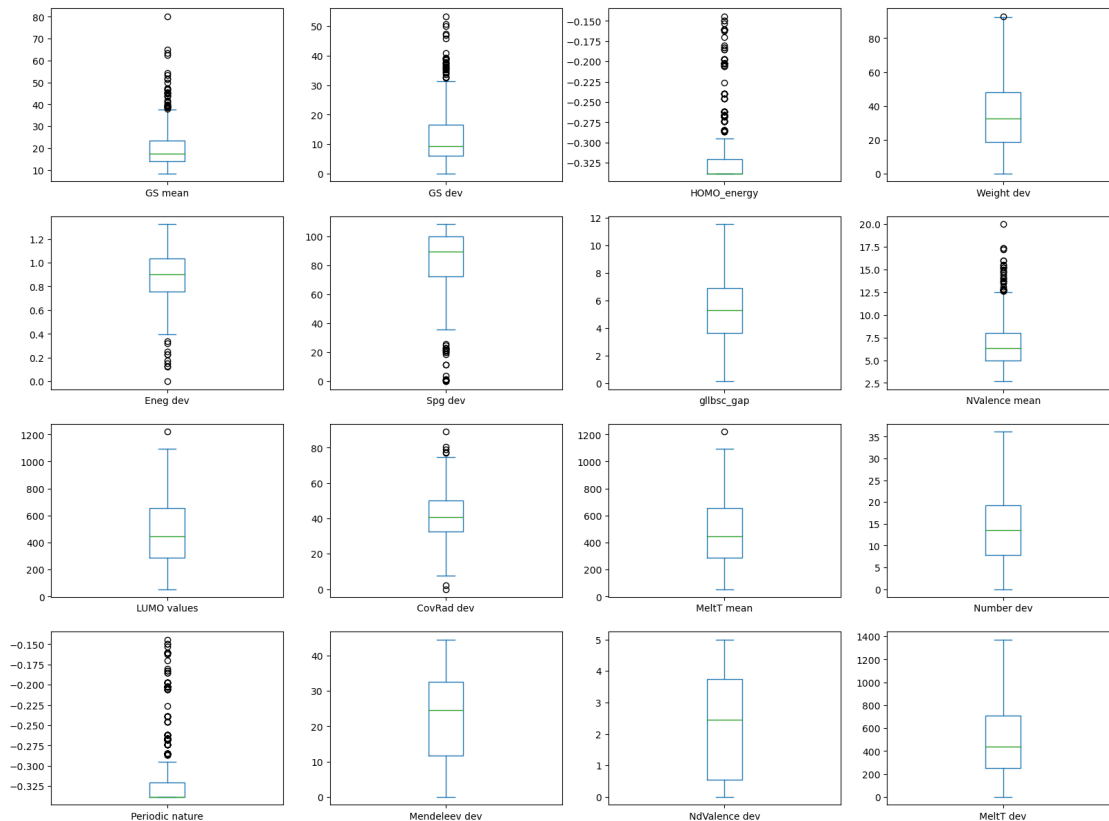
	Periodic nature	Mendeleev dev	NdValence dev	MeltT dev
16	-0.338381	6.72	1.8	736.6272

```
[24]: df = df.drop(16)
```

```
[26]: #Saving the dataset after dropping the outliers
df.to_pickle('coursework-cleaned-data.pickle')
```

```
[27]: #Plotting box plots again to make sure all outliers have been removed

columns = ['GS mean', 'GS dev', 'HOMO_energy', 'Weight dev', 'Eneg dev', 'Spg dev',
           'gllbsc_gap', 'NValence mean', 'LUMO values', 'CovRad dev',
           'MeltT mean', 'Number dev', 'Periodic nature', 'Mendeleev dev',
           'NdValence dev', 'MeltT dev']
df[columns].plot(kind='box',subplots=True,layout=(4,4),figsize=(20,15))
plt.show()
```




[28]: *#Exploring Correlations in the data*

[32]: *#Now I will obtain Pearsons correlations and make a heatmap. Doing so will*
↪allow me to look for any redundant columns

```
corrmat = df.corr()
corrmat

hm = sns.heatmap(corrmat,
                  cbar=True,
                  annot=True,
                  square=True,
                  fmt='.2f',
                  annot_kws={'size': 12},
                  yticklabels=df.columns,
                  xticklabels=df.columns,
                  cmap="Spectral_r")

plt.show()
```



```
[34]: #Linear Regression
```

```
[35]: #First, I will set up x and y and then save the data
x = df.loc[:, df.columns != "gllbse_gap"].values
y = df.loc[:, df.columns == "gllbse_gap"].values
df.to_pickle('coursework-regression-train.pickle')
```

```
[36]: #I will now look at the shape of the dataset
df.values.shape
```

```
[36]: (746, 16)
```

```
[37]: #The number of rows has decreased from 749 to 746 because I removed three rows
      ↳ due to the fact that they had outliers
```

```
[38]: #Scaling the data
      #I will now standardize the data using the StandardScaler function
```

```

from sklearn.preprocessing import StandardScaler

scaler_x = StandardScaler()
x = scaler_x.fit_transform(x)
scaler_y = StandardScaler()
y = scaler_y.fit_transform(y.reshape(-1, 1))

```

[39]: *#Here, I will use the train_test_split tool from scikit-learn to make an 80:20 ↪ training:test split as shown in the previous exercises.*

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state=0, ↪
↪train_size=0.8)

```

[40]: *#Setting up a linear regression and fitting the model*

```

#This code fits the data using the LinearRegression function from scikit-learn
from sklearn.linear_model import LinearRegression
# with sklearn
regr = LinearRegression()
regr.fit(x_train, y_train)

print('Intercept: \n', regr.intercept_)
print('Coefficients: \n', regr.coef_)

#Using the results I can determine which feature appears to have the greatest ↪
↪influence on the band gap

```

Intercept:

[0.00338702]

Coefficients:

[[-0.10255972 0.35950801 -0.12733959 -0.43379754 0.40560351 -0.2227987
0.15962037 0.03401533 -0.43188967 0.03401533 0.22291752 -0.12733959
-0.06211499 -0.5043008 -0.01283062]]

[41]: *#Based on this I believe that the feature with the greatest influence on the ↪*
↪band gap is feature (#) which is BLANK

[49]: *# Fit a baseline linear regression model*

```

lr = LinearRegression()
lr.fit(x_train, y_train)
y_pred_lr = lr.predict(x_test)
mse_lr = mean_squared_error(y_test, y_pred_lr)
print("Baseline Linear Regression MSE: ", mse_lr)

```


Baseline Linear Regression MSE: 0.4836761609819223

```
[43]: #Analysing the performance of my model

#With the below code I am first making predictions for the training set. Then I
↳ am creating a scatter plot graph to plot the predictions I made.

predictions = regr.predict(x_train)
plt.scatter(y_train, predictions)
plt.xlabel('True value')
plt.ylabel('Predicted value')
```


```
[43]: Text(0, 0.5, 'Predicted value')
```



```
[44]: #From there, I can plot the test set

predictions = regr.predict(x_test)
plt.scatter(y_test, predictions)
plt.xlabel('True value')
plt.ylabel('Predicted value')
```

```
[44]: Text(0, 0.5, 'Predicted value')
```



```
[45]: #Next, I will summarize the performance of my model using statistics values  
#These include mean squared error, root mean squared error and r-squared
```

```
from sklearn.metrics import mean_squared_error, r2_score  
  
print('Mean squared error:', mean_squared_error(predictions, y_test))  
print('Root mean squared error:', mean_squared_error(predictions, y_test,   
    ↪squared=False))  
print('r-squared:', r2_score(y_test, predictions))
```

```
Mean squared error: 0.4836761609819223  
Root mean squared error: 0.6954683033625058  
r-squared: 0.5519131120239659
```

```
[51]: #Gradient Boosted Regressor
```

```
[52]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[53]: regr = GradientBoostingRegressor()  
      regr.fit(x_train, y_train)
```


```
/opt/conda/lib/python3.10/site-packages/sklearn/ensemble/_gb.py:570:  
DataConversionWarning: A column-vector y was passed when a 1d array was  
expected. Please change the shape of y to (n_samples, ), for example using  
ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
[53]: GradientBoostingRegressor()
```

```
[54]: predictions = regr.predict(x_test)  
plt.scatter(y_test, predictions)  
plt.xlabel('True values')  
plt.ylabel('Predictions')
```

```
[54]: Text(0, 0.5, 'Predictions')
```



```
[57]: #Testing my model's performance  
  
from sklearn.metrics import mean_squared_error, r2_score  
  
print('Mean squared error:', mean_squared_error(predictions, y_test))
```



```
print('Root mean squared error:', mean_squared_error(predictions, y_test,
↪squared=False))
print('r-squared:', r2_score(y_test, predictions))
```

Mean squared error: 0.36589473580019977

Root mean squared error: 0.6048923340564003

r-squared: 0.6610280871426852

[]: